IS624 - Assignment 1

James Quacinella

06/13/2015

Code

The official R code for this homework is located on github.com. All the code below is copied from the problem functions in that file.
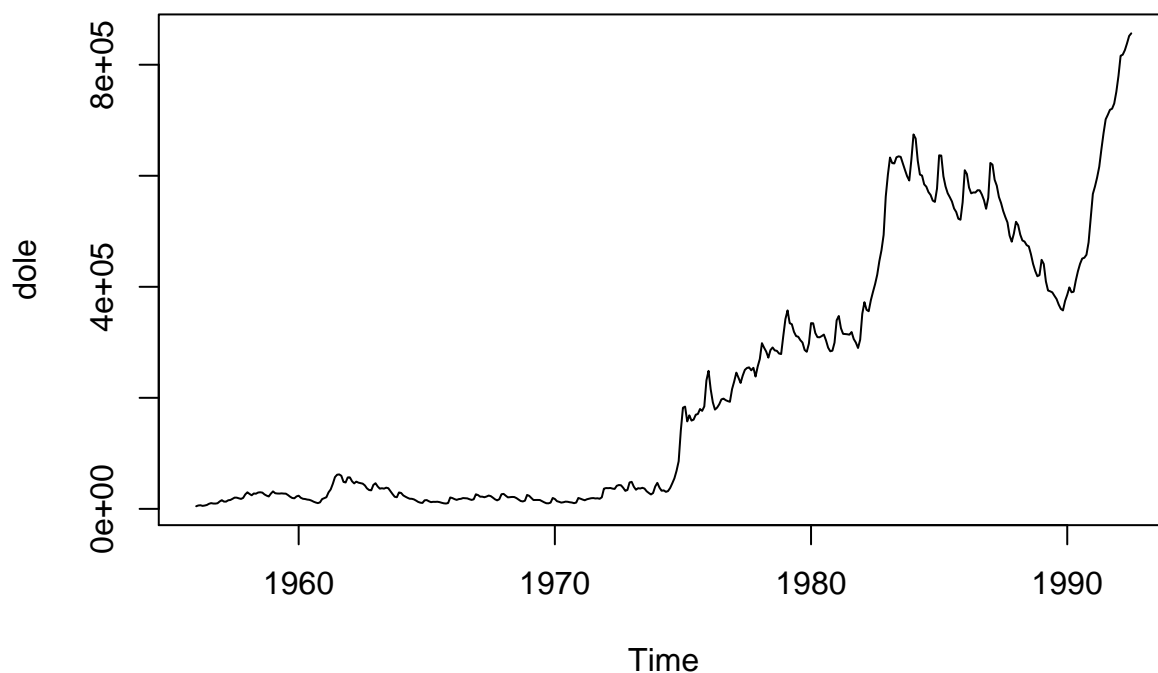
Setup
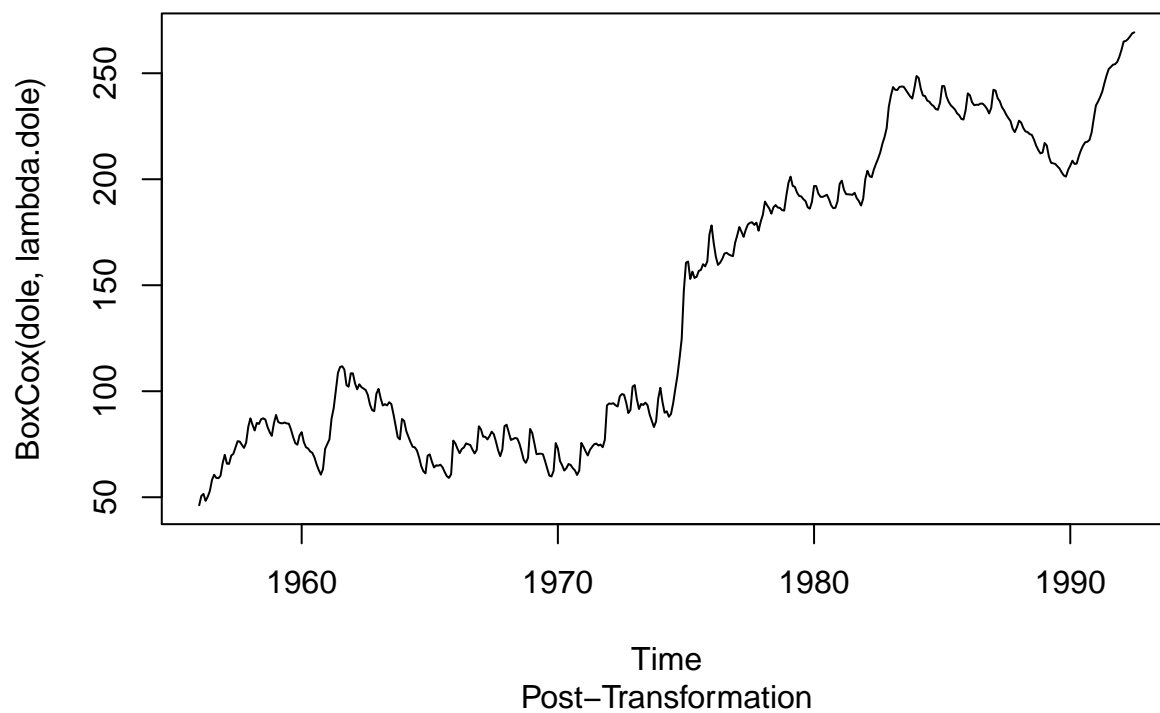
Question 2.1

Question 2.1a)

The following code produces two graphs of the pre and post transformation of the monthly total of people on unemployed benefits in Australia. TODO

```
#' Plot the Monthly total of people on unemployed benefits in Australia (January 1956-July 1992).
lambda.dole <- BoxCox.lambda(dole)   # -0.03363775
plot(dole, main="Monthly total of people on unemployed benefits in Australia", sub="Pre-Transformation")
plot(BoxCox(dole, lambda.dole), main="Monthly total of people on unemployed benefits in Australia", sub=
```

Monthly total of people on unemployed benefits in Australia

Time

Pre−Transformation

Monthly total of people on unemployed benefits in Australia
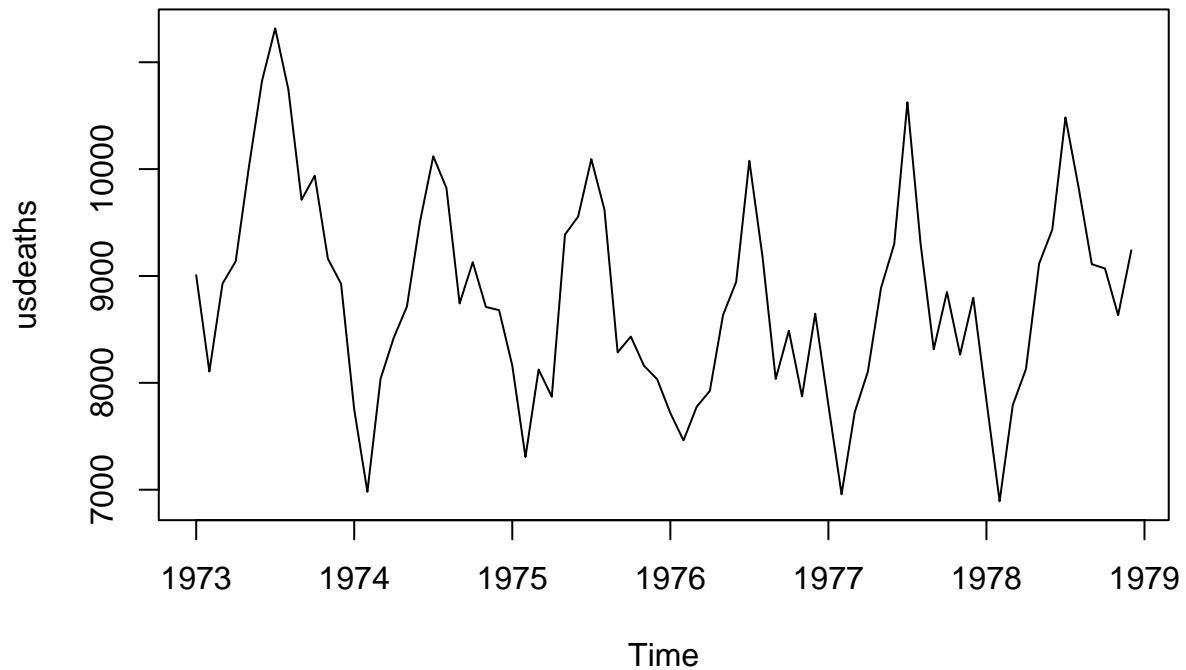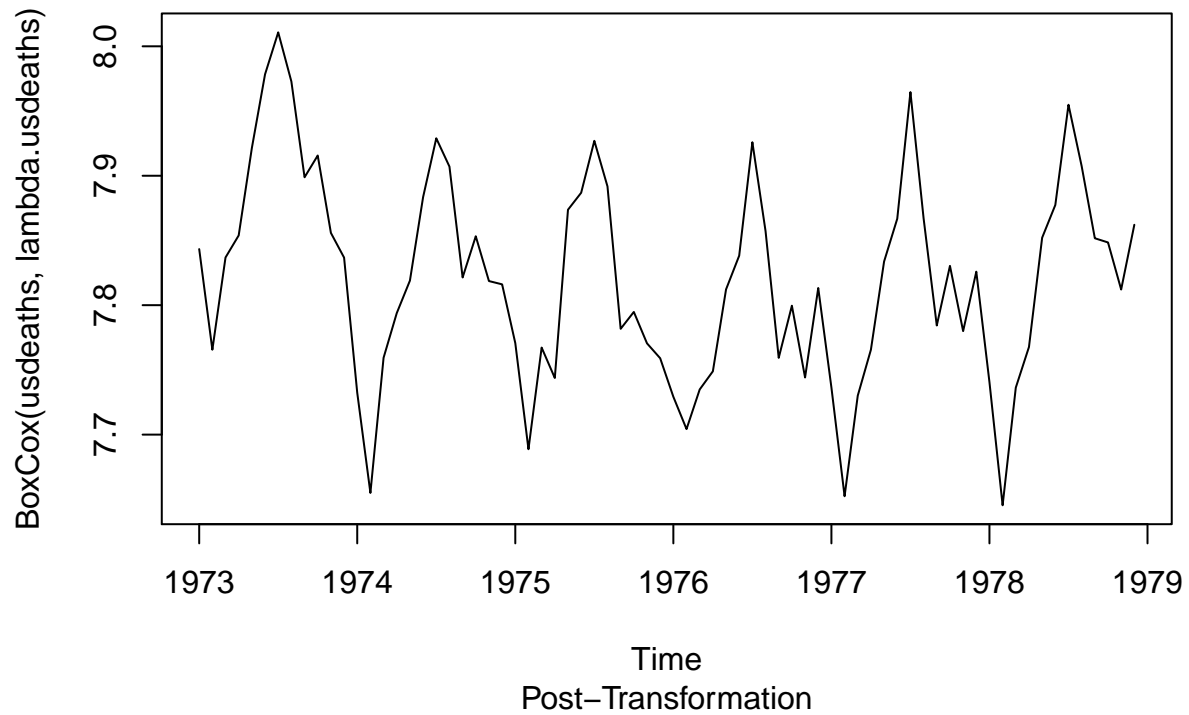

Time

Post−Transformation

Question 2.1b)

The following code produces two graphs of the pre and post transformation of the monthly total of accidental deaths in the United States. TODO

```r
#' Monthly total of accidental deaths in the United States (January 1973-December 1978).
#' The lambda value is small, but since the plot shows seasonality, I think this transformation
#' ise useful. Also, the scale post-transformaion is a lot smaller, which is sometimes helpful
#' in machine learning / predicitive tasks
lambda.usdeaths <- BoxCox.lambda(usdeaths)   # -0.03363775
plot(usdeaths, main="Monthly total of accidental deaths in the United States", sub="Pre-Transformation"
plot(BoxCox(usdeaths, lambda.usdeaths), main="Monthly total of accidental deaths in the United States",
```

Monthly total of accidental deaths in the United States

Time
Pre−Transformation

Monthly total of accidental deaths in the United States
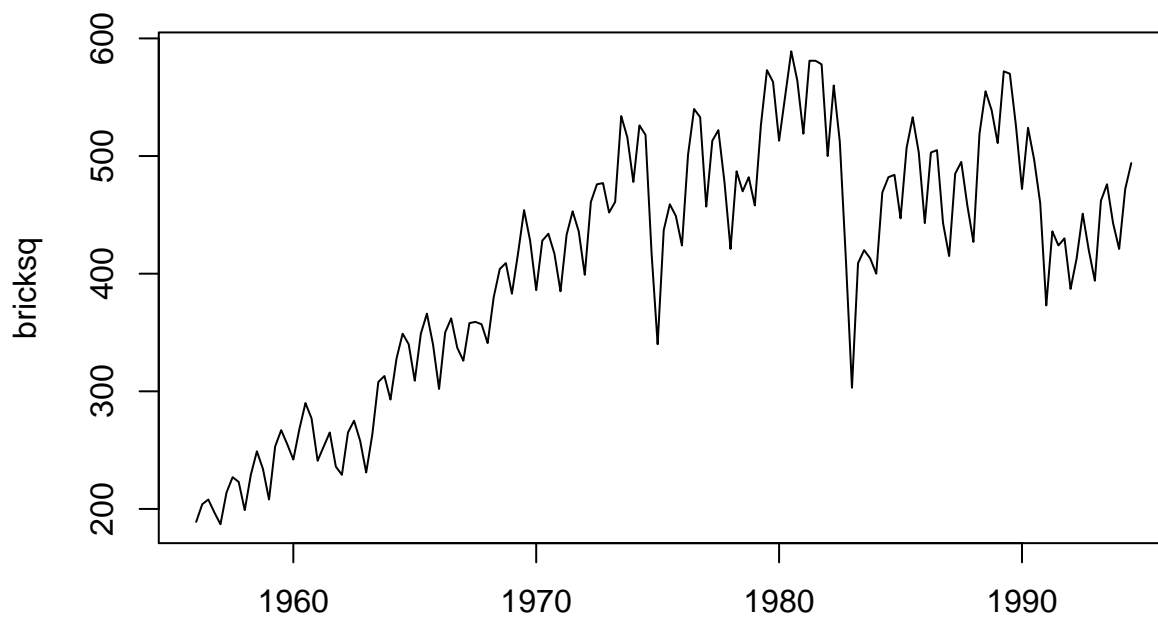

Time
Post−Transformation

Question 2.1c)

The following code produces two graphs of the pre and post transformation of the Quarterly production of bricks. TODO

```r
#' Quarterly production of bricks (in millions of units) at Portland, Australia (March 1956-September 1.
lambda.bricksq <- BoxCox.lambda(bricksq)   # 0.2548929
plot(bricksq, main="Quarterly production of bricks", sub="Pre-transformation")
plot(BoxCox(bricksq, lambda.bricksq), main="Quarterly production of bricks", sub="Post-transformation")
```

Quarterly production of bricks

Pre−transformation

Quarterly production of bricks


Post−transformation
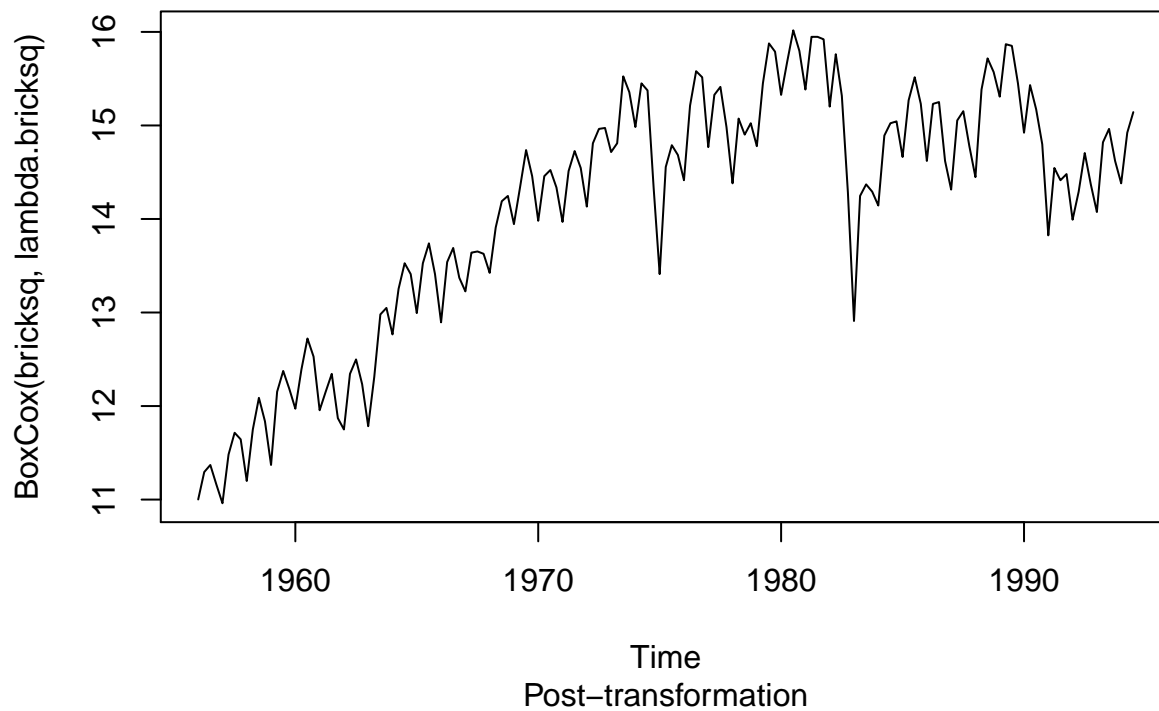
Question 2.3

Question 2.3a)

To become familiar with the ibmclose data set, here we plot the time series with some running averages. The big dip in price is going to give us trouble with our simplistic forcasting.

```r
#' Produce some plots of the data in order to become familiar with it.
plot(ibmclose, main="Closing Price of IBM Stock")
lines(rollmean(ibmclose, k=30, fill=NA), col="red")
lines(rollmean(ibmclose, k=60, fill=NA), col="blue")
legend("topright", lty=1, col=c("red", "blue"), legend=c("30 day rolling average","60 day rolling averag
```

Closing Price of IBM Stock
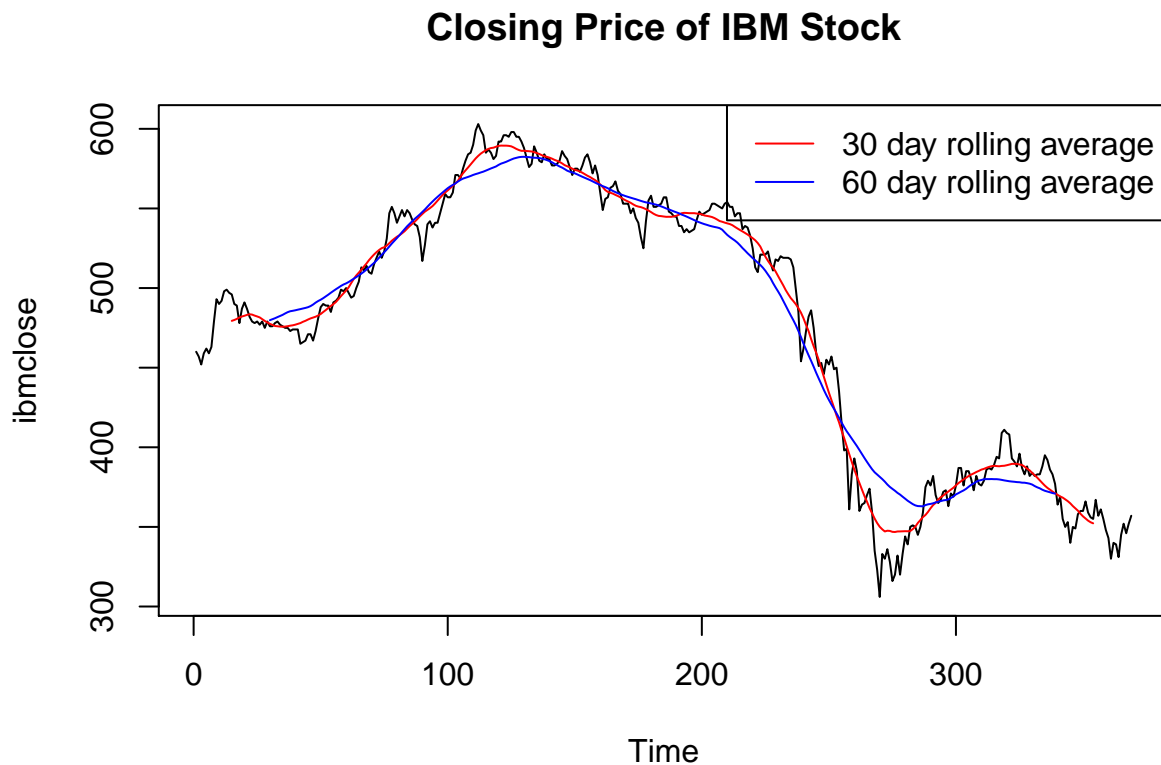
Question 2.3b)

The ibmclose data is split into a 300-day training set and a 69-day test set:

```r
ibmclose.training <- window(ibmclose, start=1, end=300)
ibmclose.test <- window(ibmclose, start=301)
```

Question 2.3c)

Note that the blue prediction is for the mean predicition, and looks different due to using plot() versus lines(). Also note that both naive methods overlap their predicitions. Visually inspecting the predictions, drift is clearly the most accurate; the mean is so far away as it does not account well for the bug dip in value around t=250.

```
ibmclose.prediction_horizon <- length(ibmclose) - 300;

# Calc mean, naive, seasonal naive, drift predictions
ibmclose.fit.mean <- meanf(ibmclose.training, h=ibmclose.prediction_horizon)
ibmclose.fit.naive <- naive(ibmclose.training, h=ibmclose.prediction_horizon)
ibmclose.fit.seasonalNaive <- snaive(ibmclose.training, h=ibmclose.prediction_horizon)
ibmclose.fit.drift <- rwf(ibmclose.training, drift=TRUE, h=ibmclose.prediction_horizon)

# Plot data with predictions
#plot(ibmclose.training, main="Closing Price of IBM Stock")
plot(ibmclose.fit.mean, plot.conf=FALSE, main="Closing Price of IBM Stock")
lines(ibmclose.fit.naive$mean, col=2)
lines(ibmclose.fit.seasonalNaive$mean, col=3)
lines(ibmclose.fit.drift$mean, col=1)
lines(ibmclose)
legend("bottomleft", lty=1, col=c(4,2,3,1), legend=c("Mean method","Naive method","Seasonal naive metho
```

Closing Price of IBM Stock
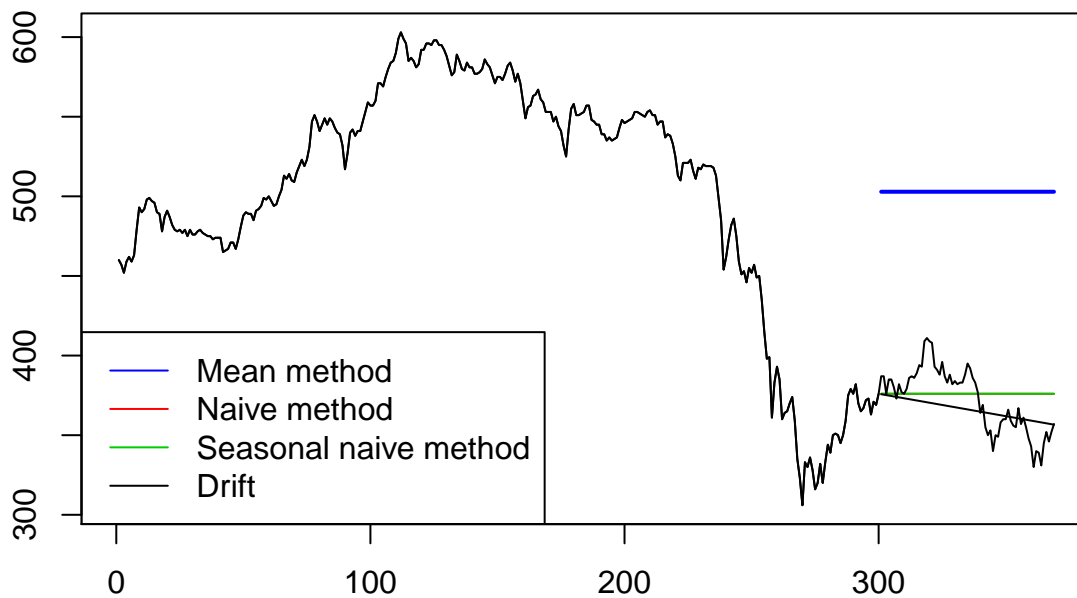
The error rates are also printed for the various prediction methods: TODO

```
# Look at error values from our predictions
accuracy(ibmclose.fit.mean, ibmclose)
accuracy(ibmclose.fit.naive, ibmclose)
accuracy(ibmclose.fit.seasonalNaive, ibmclose)
accuracy(ibmclose.fit.drift, ibmclose)
```

```
##                        ME       RMSE        MAE        MPE      MAPE
## Training set  1.660438e-14   73.61532   58.72231   -2.642058  13.03019
## Test set     -1.306180e+02  132.12557  130.61797  -35.478819  35.47882
##                     MASE       ACF1  Theil's U
## Training set  11.52098  0.9895779         NA
## Test set      25.62649  0.9314689   19.05515
```

```
##                     ME      RMSE      MAE        MPE      MAPE      MASE
## Training set -0.2809365   7.302815   5.09699 -0.08262872 1.115844 1.000000
## Test set     -3.7246377 20.248099 17.02899 -1.29391743 4.668186 3.340989
##                   ACF1 Theil's U
## Training set 0.1351052        NA
## Test set     0.9314689  2.973486
##                     ME      RMSE      MAE        MPE      MAPE      MASE
## Training set -0.2809365   7.302815   5.09699 -0.08262872 1.115844 1.000000
## Test set     -3.7246377 20.248099 17.02899 -1.29391743 4.668186 3.340989
##                   ACF1 Theil's U
## Training set 0.1351052        NA
## Test set     0.9314689  2.973486
##                       ME      RMSE      MAE         MPE      MAPE
## Training set 2.870480e-14   7.297409   5.127996 -0.02530123 1.121650
## Test set     6.108138e+00 17.066963 13.974747  1.41920066 3.707888
##                   MASE      ACF1 Theil's U
## Training set 1.006083 0.1351052        NA
## Test set     2.741765 0.9045875  2.361092
```

Question 2.4

Question 2.4a)

Here a simple ts plot, with a few moving averages is plotted, as well as a month and season plot, since the data is the right format (versus previous question).

```r
# Plot time series with some moving averages
plot(hsales, main="Sales of new one-family houses in the USA")
lines(rollmean(hsales, k=30, fill=NA), col="red")
lines(rollmean(hsales, k=60, fill=NA), col="blue")

# Show seasonal plot, which should a general downward trend from March onwards
seasonplot(hsales, col=1:20, pch=19, year.labels=TRUE, year.labels.left=TRUE)
```

Sales of new one–family houses in the USA

Seasonal plot: hsales
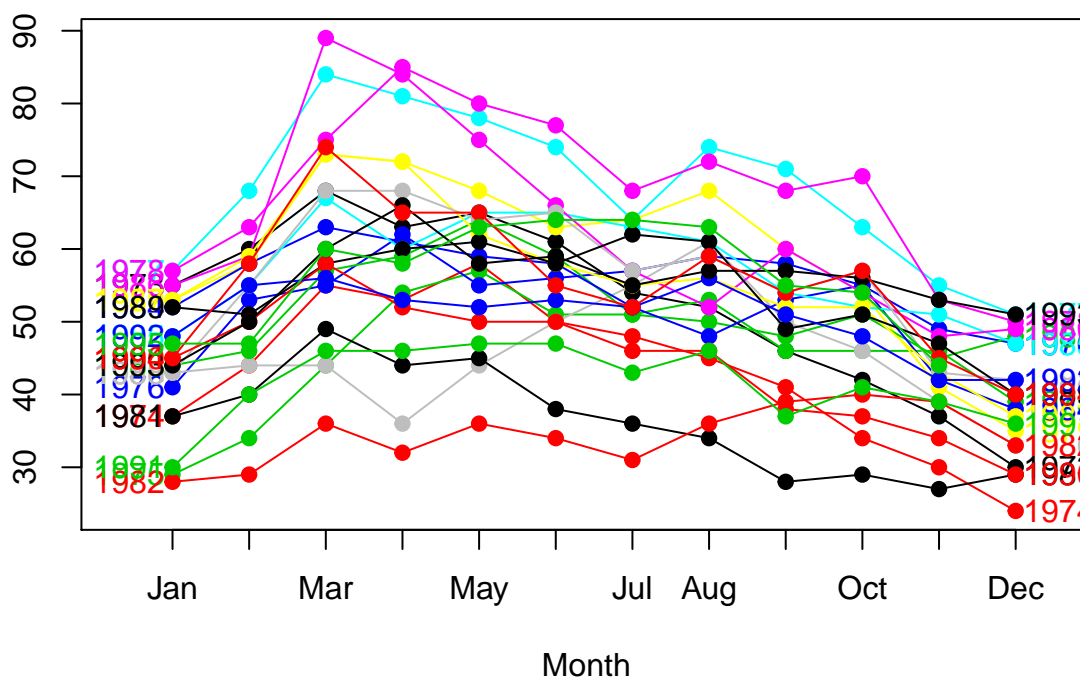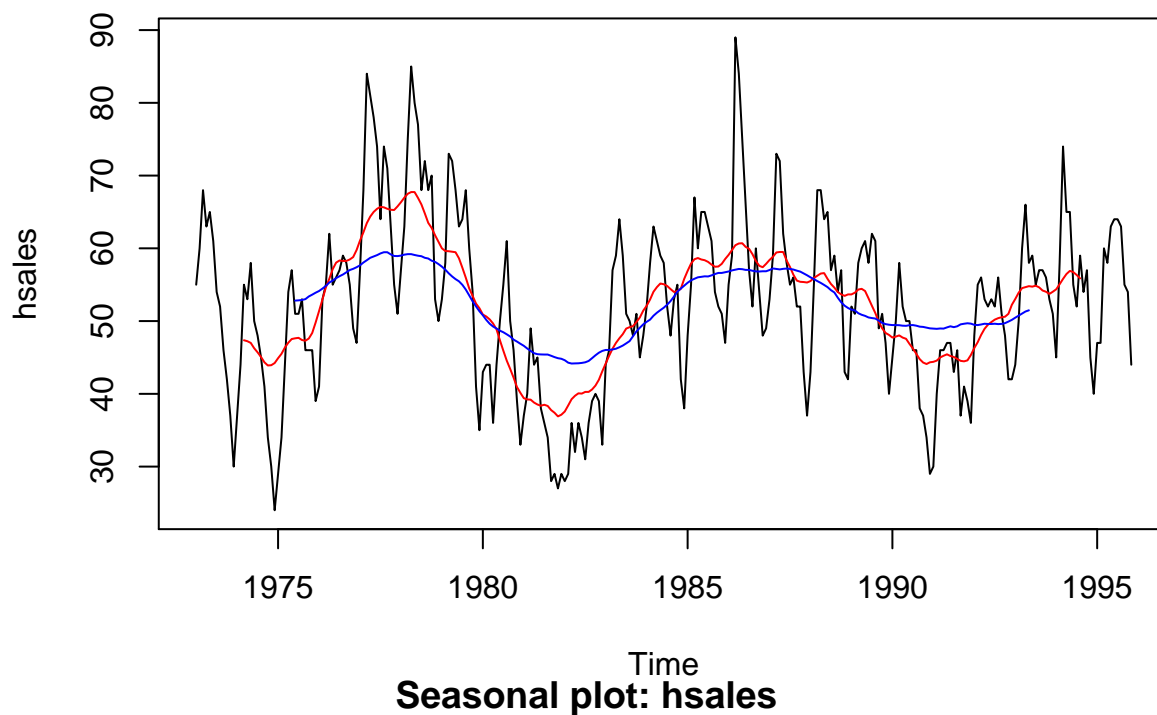
Here are the sales by month, which emphasizes that there is a ramp up in sales from Jan to Mar, but declines for the rest of the year:
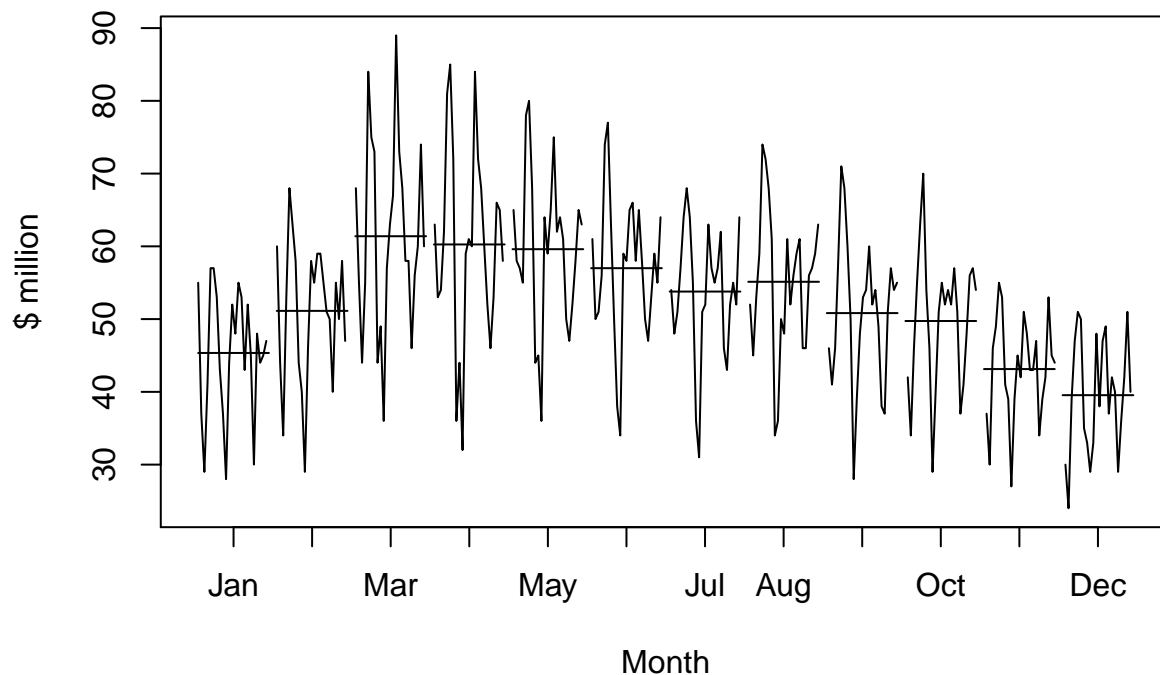
```
monthplot(hsales,ylab="$ million",xlab="Month",xaxt="n", main="Sales of new one-family houses in the USA
axis(1,at=1:12,labels=month.abb,cex=0.8)
```

Sales of new one–family houses in the USA

Question 2.4b)

Split the hsales data set into a training set and a test set, where the test set is the last two years of data.

```
hsales.training <- window(hsales, start=c(1973, 1), end=c(1993, 12))
hsales.test <- window(hsales, start=c(1994, 1), end=c(1995, 11))
```

Question 2.4c)

Visual inspection shows that the seasonal naive method looks pretty accurate across the 23months of prediction, though the predicted value in 23months looks pretty off. None of the methods seem to predict the end drop-off in value towards the end of 1995. TODO

```
hsales.prediction_horizon <- 23;

# Calculate the mean, naive, seasonal naive, drift predicitons
hsales.fit.mean <- meanf(hsales.training, h=hsales.prediction_horizon)
hsales.fit.naive <- naive(hsales.training, h=hsales.prediction_horizon)
hsales.fit.seasonalNaive <- snaive(hsales.training, h=hsales.prediction_horizon)
hsales.fit.drift <- rwf(hsales.training, drift=TRUE, h=hsales.prediction_horizon)

# Plot the data with our predicitons
#plot(hsales.training, main="Sales of new one-family houses in the USA w/ Forecasts")
plot(hsales.fit.mean, plot.conf=FALSE, main="Sales of new one-family houses in the USA")
lines(hsales.fit.naive$mean, col=2)
lines(hsales.fit.seasonalNaive$mean, col=3)
lines(hsales.fit.drift$mean, col=1)
```

```
lines(hsales)
legend("topright", lty=1, col=c(4,2,3,1), legend=c("Mean method","Naive method","Seasonal naive method"
```

Sales of new one–family houses in the USA
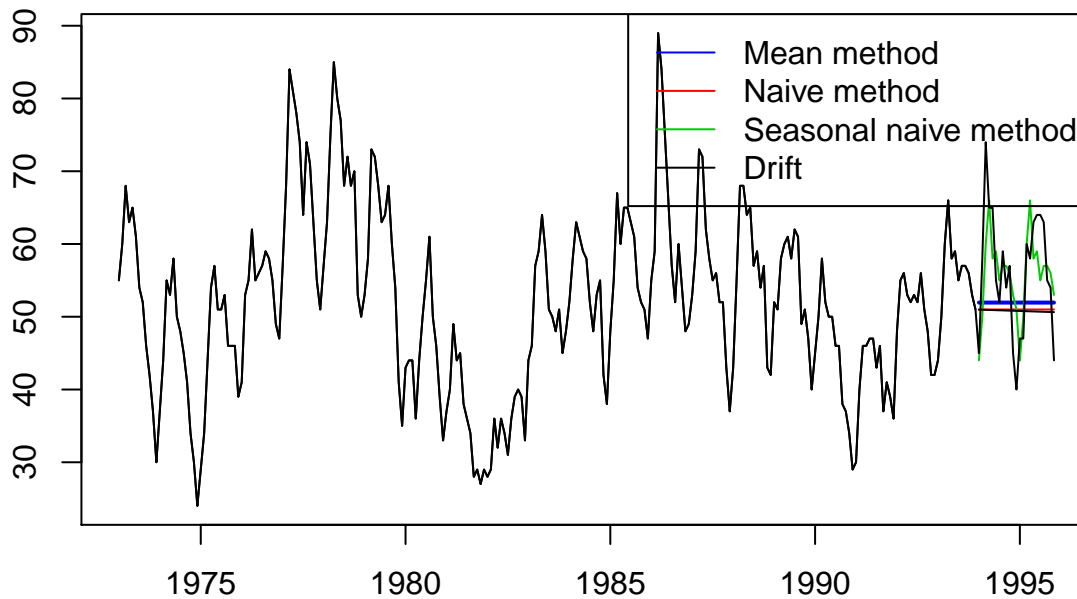
Here we look at the error values (TODO):

```
# Look at error values from our predictions
accuracy(hsales.fit.mean, hsales)
accuracy(hsales.fit.naive, hsales)
accuracy(hsales.fit.seasonalNaive, hsales)
accuracy(hsales.fit.drift, hsales)
```

```
##                         ME      RMSE      MAE        MPE      MAPE     MASE
## Training set 2.480763e-15 12.138802 9.498898 -6.120182 20.30851 1.119163
## Test set     4.051587e+00  9.216133 7.850759  5.074990 13.75973 0.924979
##                   ACF1 Theil's U
## Training set 0.8661515        NA
## Test set     0.5095178   1.13105
##                        ME      RMSE      MAE        MPE      MAPE      MASE
## Training set -0.01593625  6.289813 4.988048 -0.7800232  9.880157 0.5876934
## Test set      5.00000000  9.670664 8.304348  6.8080182 14.381673 0.9784210
##                   ACF1 Theil's U
## Training set 0.1829708        NA
## Test set     0.5095178  1.179633
##                      ME      RMSE    MAE        MPE     MAPE      MASE
## Training set 0.1375000 10.576113 8.4875 -2.1016380 17.63375 1.0000000
## Test set     0.3043478  6.160886 5.0000 -0.7312374  9.12828 0.5891016
##                  ACF1 Theil's U
## Training set 0.838108        NA
## Test set     0.224307 0.8031005
##                        ME      RMSE      MAE        MPE     MAPE      MASE
## Training set 2.377739e-15 6.289793 4.987730 -0.7474544  9.87819 0.5876560
```

```
## Test set      5.191235e+00 9.761548 8.393037  7.1599507 14.50303 0.9888703
##                     ACF1 Theil's U
## Training set 0.1829708        NA
## Test set     0.5083059  1.188562
```
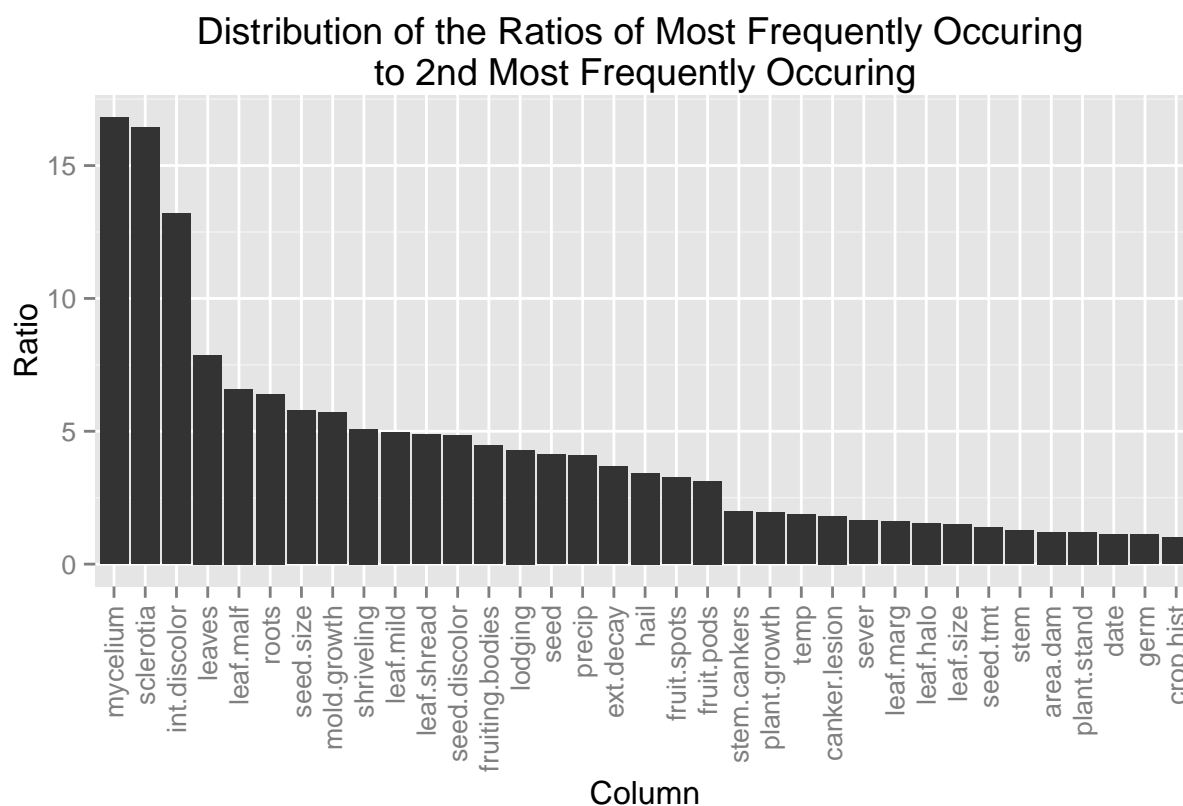
Question 3.2a

It seems that only three columns come close (but don't quite hit the threshold) to matching the rules of thumb for degenerate categorical data as described in the book. Those columns are mycelium, int.discolor, and sclerotia (see output):

```
nrows <- nrow(Soybean);
i <- 1;
ratios <- c()
for (column in colnames(Soybean)[2:length(colnames(Soybean))]) {
  t <- count(Soybean, column)
  sorted_t <- t[order(-t['freq']), 'freq']

  ratios[i] <- (sorted_t[1] / sorted_t[2]);
  i <- i + 1;
}

ratios.dist <- data.frame(y=ratios, x=colnames(Soybean)[2:length(colnames(Soybean))])
ratios.dist$x <-reorder(ratios.dist$x,-ratios.dist$y)
ggplot(ratios.dist) +
  geom_bar(aes(x=x, y=y), stat="identity") +
  theme(axis.text.x=element_text(angle=90,hjust=1,vjust=0.5)) +
  ggtitle("Distribution of the Ratios of Most Frequently Occuring \nto 2nd Most Frequently Occuring") +
  ylab("Ratio") + xlab("Column")
```


15

Question 3.2b

Looking at the histogram, it seems that quite a few predictors have signifcant number of NAs. There is a dropoff in NA level from around 12% to 5% (highest being almost 20% missing data). Also, the columns herbicide-injury, cyst-nematode, 2-4-d-injury, diaporthe-pod-&-stem-blight have no instances / rows that do not have a NA in them. The class phytophthora-rot has 68 NAs, but all other classes have complete cases:

```r
i <- 1;
nas <- c()
for (column in colnames(Soybean)[2:length(colnames(Soybean))]) {
  t <- count(Soybean, column)
  nas[i] <- ifelse(length(t[is.na(t[column]), 'freq']), t[is.na(t[column]), 'freq'] * 100 / nrows, 0)
  i <- i + 1;
}

# Display histogram of distrubtion of NAs
nas.dist <- data.frame(y=nas, x=colnames(Soybean)[2:length(colnames(Soybean))])
nas.dist$x <-reorder(nas.dist$x,-nas.dist$y)
ggplot(nas.dist) +
  geom_bar(aes(x=x, y=y), stat="identity") +
  theme(axis.text.x=element_text(angle=90,hjust=1,vjust=0.5)) +
  ggtitle("Distribution of Missing Data") +
  ylab("% Rows w/ NAs") + xlab("Column")
```
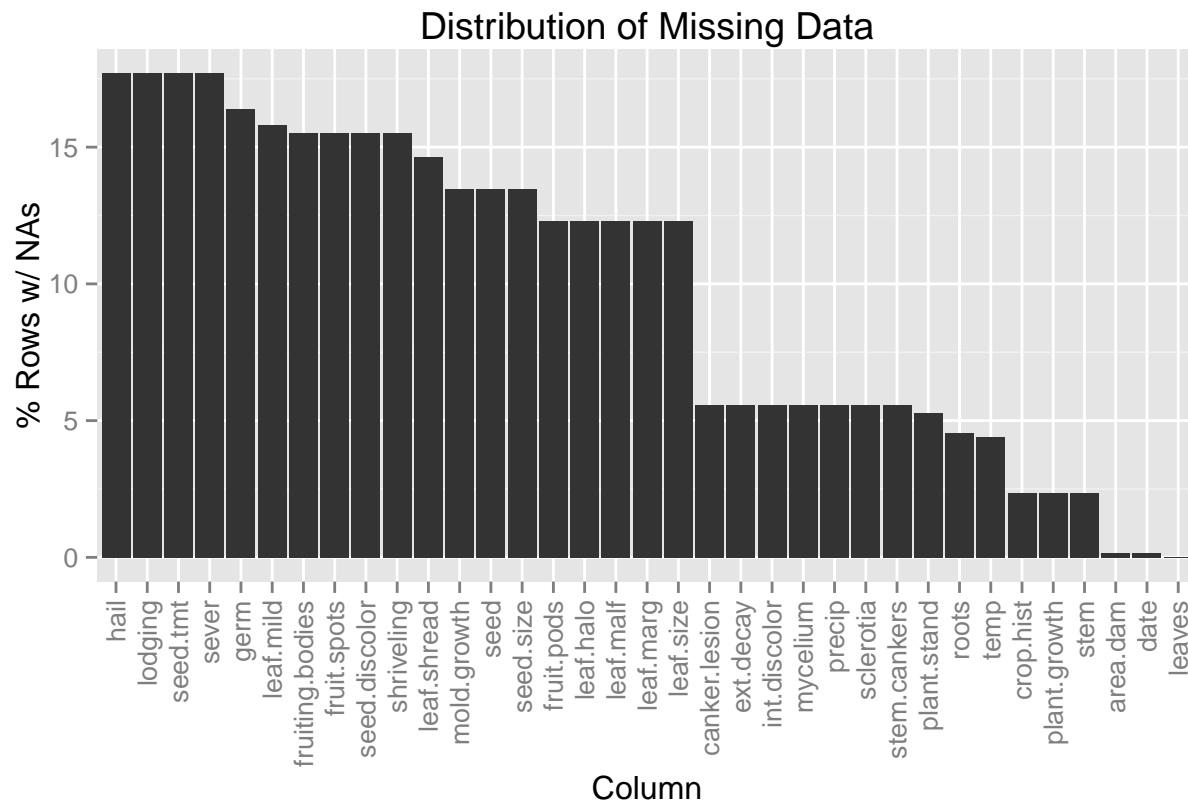


```r
# The columns herbicide-injury, cyst-nematode, 2-4-d-injury, diaporthe-pod-&-stem-blight
# have no instances / rows that do not have a NA in them
# The Class phytophthora-rot has 68  NAs, but all other classes have none
withnas <- count(Soybean$Class)
```

```
withoutnas <- count(Soybean[complete.cases(Soybean), 'Class'])
for (i in 1:nrow(withoutnas)) {
  classval <- withoutnas[i, 'x']
  diff <- withnas[withnas$x == classval, 'freq'] - withoutnas[withoutnas$x == classval, 'freq']
  if(diff > 0) {
    print(paste("Class", classval, "has", diff, " NAs"))
  }
}
```

```
## [1] "Class phytophthora-rot has 68  NAs"
```

Question 3.2c)

We should start by investigating the nature of the missing data for predictors with a large number of missing values as well as why some classes have a high occurance of NAs. Our predictive model is going to be biased if we do not have good observations for 4 or 5 of the 19 outcome classes.

For each predictor that has a high number of NAs (starting from the left in the question_3.2b histogram), we can determine if we can remove the predictor by checking if it has a high collinearity with any other predictor. If so, we can remove that preidctor and simply use the collinear predictor (presuming it has less instances of NAs). Imputing values would be difficult, since these are all categorical variables, and averaging those values would make no sense. We could impute a value using K-nearest neighbors, which would find the K rows that are closest to it by some definition of 'close', and then take a majority vote on what the categorical value should be.

Question 4.4a)

TODO

```
#' (a) Use the sample function in base R to create a completely random sample
#' of 60 oils. How closely do the frequencies of the random sample match
#' the original samples? Repeat this procedure several times of understand
#' the variation in the sampling process.
print("Frequency from Data")
print(freq(oilType, plot = FALSE))

cat("\n")
print("Frequency from Sampling 60 items")
for (i in 1:5) {
  print(freq(sample(oilType, 60), plot = FALSE))
}
```

```
## [1] "Frequency from Data"
## oilType
##        Frequency Percent
## A             37  38.542
## B             26  27.083
## C              3   3.125
## D              7   7.292
## E             11  11.458
## F             10  10.417
## G              2   2.083
## Total         96 100.000
##
## [1] "Frequency from Sampling 60 items"
## sample(oilType, 60)
##        Frequency Percent
## A             21  35.000
## B             19  31.667
## C              1   1.667
## D              6  10.000
## E              5   8.333
## F              7  11.667
## G              1   1.667
## Total         60 100.000
## sample(oilType, 60)
##        Frequency Percent
## A             21   35.00
## B             16   26.67
## C              3    5.00
## D              6   10.00
## E              7   11.67
## F              7   11.67
## G              0    0.00
## Total         60  100.00
## sample(oilType, 60)
##        Frequency Percent
## A             21  35.000
## B             17  28.333
```

```
## C                 2   3.333
## D                 5   8.333
## E                 9  15.000
## F                 5   8.333
## G                 1   1.667
## Total            60 100.000
## sample(oilType, 60)
##         Frequency Percent
## A              25  41.667
## B              14  23.333
## C               3   5.000
## D               4   6.667
## E               7  11.667
## F               6  10.000
## G               1   1.667
## Total          60 100.000
## sample(oilType, 60)
##         Frequency Percent
## A              25  41.667
## B              18  30.000
## C               0   0.000
## D               4   6.667
## E               6  10.000
## F               6  10.000
## G               1   1.667
## Total          60 100.000
```

Question 4.4b)

The createDataPartition is indeed giving back different numbers every time but table() output shows how it is very consistent in how it distributes samples across class values, and is very close to the ideal from the data itself. Much better than the previous approach.

```r
#' (b) Use the caret package function createDataPartition to create a stratified
#' random sample. How does this compare to the completely random samples?
#' the variation in the sampling process.
print("Frequency from Data")
print(freq(oilType, plot = FALSE))

cat("\n")
print("Frequency from Sampling 60 items")
for (i in 1:5) {
  print(freq(oilType[createDataPartition(oilType, p=60 / length(oilType), list=FALSE)], plot = FALSE))
}
```

```
## [1] "Frequency from Data"
## oilType
##         Frequency Percent
## A              37  38.542
## B              26  27.083
## C               3   3.125
## D               7   7.292
## E              11  11.458
```

```
## F                 10  10.417
## G                  2   2.083
## Total             96 100.000
##
## [1] "Frequency from Sampling 60 items"
## oilType[createDataPartition(oilType, p = 60/length(oilType),      list = FALSE)]
##         Frequency Percent
## A                24  37.500
## B                17  26.562
## C                 2   3.125
## D                 5   7.812
## E                 7  10.938
## F                 7  10.938
## G                 2   3.125
## Total            64 100.000
## oilType[createDataPartition(oilType, p = 60/length(oilType),      list = FALSE)]
##         Frequency Percent
## A                24  37.500
## B                17  26.562
## C                 2   3.125
## D                 5   7.812
## E                 7  10.938
## F                 7  10.938
## G                 2   3.125
## Total            64 100.000
## oilType[createDataPartition(oilType, p = 60/length(oilType),      list = FALSE)]
##         Frequency Percent
## A                24  37.500
## B                17  26.562
## C                 2   3.125
## D                 5   7.812
## E                 7  10.938
## F                 7  10.938
## G                 2   3.125
## Total            64 100.000
## oilType[createDataPartition(oilType, p = 60/length(oilType),      list = FALSE)]
##         Frequency Percent
## A                24  37.500
## B                17  26.562
## C                 2   3.125
## D                 5   7.812
## E                 7  10.938
## F                 7  10.938
## G                 2   3.125
## Total            64 100.000
## oilType[createDataPartition(oilType, p = 60/length(oilType),      list = FALSE)]
##         Frequency Percent
## A                24  37.500
## B                17  26.562
## C                 2   3.125
## D                 5   7.812
## E                 7  10.938
## F                 7  10.938
## G                 2   3.125
```

```
## Total       64 100.000
```

Question 4.4c)

When the number of samples is not large, a single test set should be avoided because we may need every sample during model building. We should use somekind of resampling technique, like k-fold cross-validation. The book recommends "If the samples size is small, we recommend repeated 10-fold cross-validation for several reasons: the bias and variance properties are good and, given the sample size, thecomputational costs are not large.

Question 4.4d)

The code following this shows the commands I used to come to the following conclusions.

By keeping the model performance steady, while changing the test sample size leads to a relationship where higher sample size means less uncertainty. This makes sense: we would expect us to be more certain if we have more samples to evaluate the model performance with.

If you keep the sample size steady, and vary the model performance leads to a relationship where: - the uncertainity maxes when performance is 50% (which makes sense since its binomial probability). - the uncertainity is less when model performace is near either extreme (partially due to clipping of the CI by the boundry)

```
spread <- function(obj) { conf <- obj$conf.int; return(conf[2] - conf[1]);}
spread(binom.test(16, 20)) # 0.37928
spread(binom.test(32, 40)) # 0.2659556
spread(binom.test(30, 40)) # 0.2850472
spread(binom.test(10, 40)) # 0.2850472
spread(binom.test(15, 40)) # 0.3147225
spread(binom.test(20, 40)) # 0.3239644
spread(binom.test(25, 40)) # 0.3147225
spread(binom.test(40, 40)) # 0.0880973
spread(binom.test(20, 40)) # 0.3239644
spread(binom.test(21, 40)) # 0.3236002
spread(binom.test(19, 40)) # 0.3236002
spread(binom.test(19, 20)) # 0.2474677
spread(binom.test(20, 20)) # 0.1684335
spread(binom.test(10, 20)) # 0.4560843
spread(binom.test(1, 20))  # 0.2474677
```